

E-tivity 1



Gabriel Espinoza Rojas

22176691

Risk, Ethics and Governance of AI

IN5103

University of Limerick

0.1 Task 2

The myth of the golem, marked in the infographic as **ancient**, has been part of Jewish folklore dating back to the Old Testament (Rozenberg and Leviant 2007). My favourite retelling of the story is Wisniewski (1996)’s:

“A saintly rabbi miraculously brings to life a clay giant who helps him watch over the Jews of sixteenth-century Prague.”

In my opinion, an historically overlooked attribute of the golem is its vision. In fiction, the ability to see and understand the world comes as expected; the golem is a guardian, its vision is required. As a computer engineer, over the last three years, I have been involved in a road management project and I can say from first-hand experience that incident detection is not an easy task. The radars can only take you so far and now more than ever, the idea of casting silicon-made giant golems that can watch over the roads sounds less like magic and more like another AI-powered solution.

Motorways have the highest speed limits and cover the longest distances. Have the AI vision models reached the trustworthiness of the fantastical golem’s eyes? Moreover, road management is governmental and political; would nations accept help from an automaton? Or would politics push the creators to undo their golems? just like the demise of Joseph the golem, terminated by his creator after a political treaty (Wisniewski 1996). Before sharing my thoughts, I want to review the recent history around this topic.

In **2006**, the UK Parliament implemented the first trial on smart motorways on the M42. From junction 3a to junction 7 (UK Parliament 2019), inductive loop detectors and variable message signs (VMS) helped control traffic flow on the hard shoulder. A useful solution for relieving congestion, however, despite the UK Parliament calling it ‘smart’, to today’s standards, the term seems like a stretch. And no vision component was part of it. But when put into context, this same year, NVIDIA was releasing for the very first time a new parallel computing platform: Compute Unified Device Architecture, or CUDA (NVIDIA 2026).

CUDA would eventually become a complete game-changer; it will allow software developers to reach computational levels never even dreamed of before, but

not in 2006. However, another important event from this year is that Dr. Fei-Fei Li Deng, often referred as the ‘Godmother of AI’ (especially for her contributions to the computer vision discipline), was already formulating rough ideas of a new project she thought necessary.

In **2009**, ‘ImageNet: A large-scale hierarchical image database’ was formally introduced to the computer vision community (Deng et al. 2009). Dr Li and her colleagues knew that large amounts of data would be key to developing and testing vision models. And they were not wrong. Just a year later (**2010**), they were organising the first edition of ‘ImageNet Large Scale Visual Recognition Challenge’ or ILSVRC (Russakovsky et al. 2015). Initiatives like this were great for the machine learning community, bringing it focus and competitiveness. For everyone else, it probably did not mean much, but for this niche, the air became electric and something huge was bound to happen. It did not happen the next year (2011), but the year after.

In **2012**, AlexNet (Krizhevsky et al. 2017) won the ILSVRC classification task, beating the second place (the model SIFT + FV) with a margin of 11.1% less error. This shook the foundation of object classification models in a way never seen before. It was the perfect storm that had been forming over the vision horizon for the last six years. Better hardware (thanks to CUDA giving access to the GPU) and rescued ideas from the past (of mimicking the human brain) made possible nearly a 50% reduction in object detection error, something unheard of before the use of deep learning and convolutional neural networks (CNN).

In this was not a fluke either, the next (**2013**), the winners of ILSVRC also made use of CNNs to an absolute drop of 3.6 percentage points, or about a 23.5% relative reduction in error in comparison to AlexNet (Zeiler 2012). As time passed, different combinations of deep learning techniques and hardware GPU coding

In **2017**, a new architecture shifted interest towards attention (Vaswani et al. 2017)

In **2020**, a new architecture shifted interest towards attention (Carion et al. 2020)

In **2024**, with the AI Act, Europe declared smart motorways (among others) as high-risk systems:

AI systems intended to be used as safety components in the management and operation of critical digital infrastructure, road traffic, or in the supply of water, gas, heating or electricity. (European Parliament and Council of the European Union 2024, Annex III, point 2).

The AI summer and the

References

- Carion, N., Massa, F., Synnaeve, G., Usunier, N., Kirillov, A. and Zagoruyko, S. (2020), End-to-end object detection with transformers, *in* ‘European Conference on Computer Vision’, pp. 213–229. iii
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K. and Fei-Fei, L. (2009), Imagenet: A large-scale hierarchical image database, *in* ‘2009 IEEE Conference on Computer Vision and Pattern Recognition’, pp. 248–255. iii
- European Parliament and Council of the European Union (2024), ‘Regulation (eu) 2024/1689 laying down harmonised rules on artificial intelligence’. iv
- Krizhevsky, A., Sutskever, I. and Hinton, G. E. (2017), ‘Imagenet classification with deep convolutional neural networks’. iii
- NVIDIA (2026), ‘CUDA C++ Programming Guide: Introduction’, Online technical documentation. States that NVIDIA introduced CUDA in 2006 to enable computational workloads to use GPU throughput independent of graphics APIs. Accessed 25 May 2026.
URL: <https://docs.nvidia.com/cuda/cuda-programming-guide/01-introduction/introduction.html> ii
- Rozenberg, Y. Y. and Leviant, C. (2007), *The golem and the wondrous deeds of the Maharal of Prague*, 1st ed. edn, Yale University Press, New Haven. ii
- Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., Huang, Z., Karpathy, A., Khosla, A., Bernstein, M., Berg, A. C. and Fei-Fei, L. (2015), ‘Imagenet large scale visual recognition challenge’, *International journal of computer vision* **115**(3), 211–252. iii

REFERENCES

- UK Parliament (2019), ‘Written question HL17790: Smart Motorways’. Confirms the M42 J3a–7 section was the first UK smart motorway to use the hard shoulder as an additional lane at peak times in 2006. ii
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L. and Polosukhin, I. (2017), Attention is all you need, *in* ‘Advances in Neural Information Processing Systems’, Vol. 30. iii
- Wisniewski, D. (1996), *Golem*, Clarion Books, New York. ii
- Zeiler, M. D. (2012), Adadelata: An adaptive learning rate method. iii