E-tivity 3

Gabriel Espinoza Rojas
22176691

Risk, Ethics and Governance of AI

IN5103

University of Limerick

0.1 My two case studies

0.1.1 Juana the sales manager

Juana is a sales manager for a smart grid management company. Among her tasks, she puts together and sends monthly customer reports. Recently a peer left the company and now the reports Juana must deliver nearly doubled. Her performance metrics are highly dependent on customer satisfaction and clients will quickly question any delays with some prone to threaten to terminate their contracts if the reports are not sent in time, losing any clients would impact negatively both the company and her. HR explained that replacing the coworker would take a few months between recruitment, training, and onboarding.Pressured by internal and external factors, she turned to an AI service that can ingest client-site data and with examples she provides, generate the reports. At first, she is hesitant, redacting any sensitive information that could infringe GDPR, but after seeing how convincing the AI reports look and due to time constraints, she decides to submit raw information for each client that was previously handled by her colleague and let the chatbot decide the how to create each report; after all, she only needs to use the tool for a few months. A couple of weeks go by, and business goes as usual, until one day Juana checks her inbox and sees an email from HR addressing all members of the organisation reminding employees to be respectful towards customers and to avoid sending AI generated content to them. The email is not addressing any employee in particular; it reads more like a general warning. However, later that day the head of the sales department asks her to come by for a chat. As she makes her way to the office she wonders if she's in trouble. She checked all the reports before sending them, could she have missed something the customers didn't?

0.1.2 Paco the web developer

Paco has been working in the software development industry for the last 15 years. He joined a startup a decade ago and built-up seniority, being now in charge of the development team. The company is small and business deals go through. However, almost every time, Paco is asked to put together some sort of demo to show the prospects a potential software solution that they could developed for them. He recently became aware of the technological advancements in the machine learning community and decided to try an agent coding assistant to assess for himself how useful they can actually be. After a few afternoons he is convinced that the assistant will help him increase his productivity exponentially. Before, the pursuit of a new client often meant that he would have to park maintenance and the development of new features for solutions already in place; now he can pass on the any demo requirements to his agent and leave running in the background while he continues with his scheduled work. Upper management is delighted with his performance and him and the CEO are now encouraging all developers to use AI when possible. Could Paco's experience be biased? Or will productivity increase massively if all software developers, with different tasks and different levels of professional experience use agents too?

0.2 My analysis of Natalia's case study

Figure 1 captures the selected case study. Figure 2 shows a supportive mind map used for the CORE framework analysis:

Concepts: The core of this case study is *identity*. Our *voice* is a huge part of who we are. Unfortunately, individuals with *motor neurone diseases* (MNDs) progressively lose it, to their *families*' and their despair. *Assistive technologies* are not new, but *synthetic voice AI platforms* seek to return greater **autonomy**. Rodolfo's speech felt his own. Tomás can keep teaching and sharing his histories.

Ontology: MND patients are sponsored by the AI platform. This is a tremendous display of beneficence. However, it also creates tension between both parties; the transactional relationship is one-sided. Individuals might feel obligated to accept amends in their licences or marketing stunts in fear of losing their subscription.

Risk: There's potential for great **maleficence**. Ultimately, the risk is losing autonomy over your own cloned voice. Identity theft could come from within the company (changing the terms or outrightly keeping the voice), or scammers could steal it from publicly available material with access to the technology. The potential for monetary <u>loss</u> or emotional <u>harm</u> to a distressed family is latent.

Ethics: **Autonomy** requires more than signing terms of use; it requires continuing control over storage, deletion, licensing, incapacity and inheritance. For this initiative to truly be **just**, this must be addressed and clearly communicated to the sponsored individual and their relatives.

Case Study 1 (Societal Benefit): A Voice Returned

In June 2026, Costa Rican engineer Rodolfo Carvajal, who has lived with ALS for thirteen years and breathes through a ventilator, delivered the speech at his son's graduation through an AI avatar of his younger self, speaking in a voice rebuilt from old recordings. He told the audience he had kept the thing that matters most: his mind. The video travelled fast. Among those who watched it was Tomás, a 52-year-old teacher in Madrid diagnosed with motor neurone disease a year ago.

Tomás's speech therapist suggests voice banking before his speech deteriorates further. A commercial AI platform offers free licences to patients like him; thirty minutes of recordings produce a clone his family says is indistinguishable from his real voice. For Tomás the point is identity, not convenience. Assistive devices have existed for decades, but they replace a person's voice with a machine's. A clone preserves the warmth and intonation his children grew up with, which clinicians describe as central to dignity and connection. He can keep teaching. He can record stories for grandchildren he has not met yet.

His daughter reads the terms of use and hesitates. The voice model sits on the company's servers under a licence the company may amend. The free programme is goodwill, but also marketing for a platform whose paying customers generate synthetic voices at scale, the same capability behind the impersonation scams police now warn about. And the licence is in Tomás's name, signed while he can still speak for himself.

The family asks: who should hold the rights to his voice as the disease progresses, and after his death? Should a technology this beneficial be governed by its use, or restricted because of what else it can do?

Figure 1: Natalia's case study - reference for my analysis.

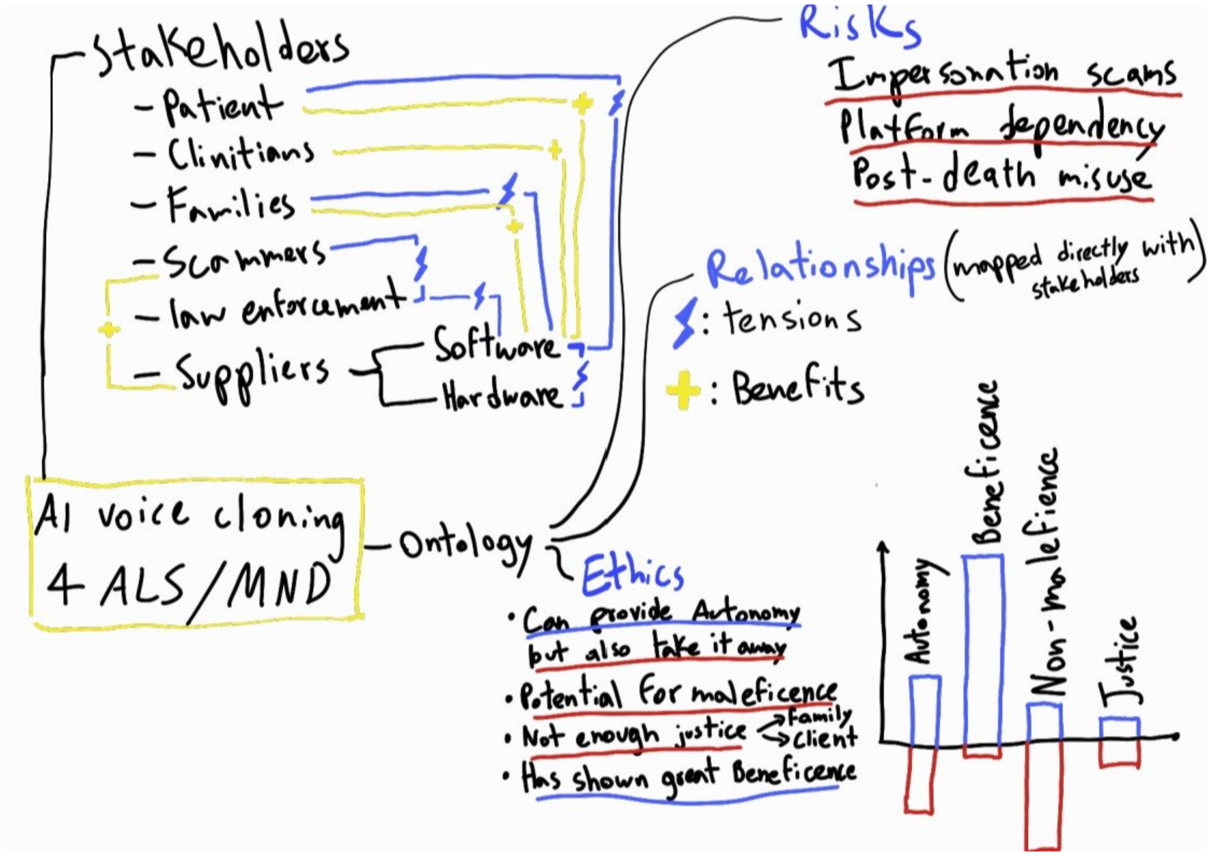

Figure 2: CORE framework mind map - Supportive diagram used for analysis.